

The Gauss Village

Himal Ghimire

June 26, 2026

1. The Fireside Tale

Before I ever set foot in that valley, I knew it.

The man who described it to me was old, his voice worn soft as saddle leather. He sat by the hearth of a roadside inn, drawing pictures in spilled ale with a crooked finger. He spoke of a village cupped between gentle hills, of a river that looped like a ribbon, of a schoolhouse under an oak, and a pond shaped precisely like a kidney bean. He named the trees: elm here, willow there, a stand of beech on the eastern slope. He estimated distances in “a stone’s throw” and “a morning’s walk.”

I wrote everything down, every detail, every number he guessed at. And that night in my room, I spread a blank sheet of paper and began to draw.

Now, understand this: I am, by trade and temperament, a cartographer. And like most cartographers of my particular school, I carry with me a tool called a **Gaussian Process**. I call it GP for short, and it is a way of building a map not just from what you have seen, but from what you *believe* about how the world behaves.

Before I drew a single contour line, I opened my field journal and wrote down the mathematics I would need. A map is a function, $f(x)$, where x is a position in the valley and $f(x)$ is the elevation (or any other feature). A Gaussian Process says that before I observe anything, any finite set of elevations follows a joint Gaussian distribution. It is completely specified by two things: a **mean function** $m(x)$ and a **covariance function** (kernel) $k(x, x')$.

For my first, naive map, I assumed the land had no particular trend, so I set the mean to zero everywhere (I could always shift it later). For the kernel, I chose the classic squared exponential, which encodes my belief in smoothness:

$$k(x, x') = \sigma_f^2 \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right) \quad (1)$$

Here σ_f^2 is the signal variance—how much the landscape varies in height overall—and ℓ is the length-scale, roughly the distance over which the elevation becomes uncorrelated. A small ℓ allows wiggly hills; a large ℓ forces gentle, rolling slopes. I set these by eye, and so my prior was:

$$f(x) \sim \mathcal{GP}(0, k(x, x')) \quad (2)$$

Then came the old man’s landmarks. I had n points: the oak at x_1 , the pond at x_2 , the schoolhouse at x_3 , and so on, each with an elevation y_i . I treated every single one as a perfect, noiseless observation: $y_i = f(x_i)$. Under the GP, the joint distribution of the

training outputs $\mathbf{y} = [y_1, \dots, y_n]^\top$ and the unknown function values $\mathbf{f}_*$ at my map grid points X_* is:

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} K(X, X) & K(X, X_*) \\ K(X_*, X) & K(X_*, X_*) \end{bmatrix} \right) \quad (3)$$

Here $K(X, X)$ is the $n \times n$ matrix of kernel evaluations between training points, $K(X, X_*)$ is $n \times n_*$, and so on. Conditioning on the observations gives the posterior predictive distribution—the actual map:

$$\mathbf{f}_* \mid X, \mathbf{y}, X_* \sim \mathcal{N}(\boldsymbol{\mu}_*, \Sigma_*) \quad (4)$$

$$\boldsymbol{\mu}_* = K(X_*, X) K(X, X)^{-1} \mathbf{y} \quad (5)$$

$$\Sigma_* = K(X_*, X_*) - K(X_*, X) K(X, X)^{-1} K(X, X_*) \quad (6)$$

These two equations were my brush and my ink. The mean $\boldsymbol{\mu}_*$ drew the contour lines. The diagonal of Σ_* gave the variance at each grid point—the pale halos of uncertainty I painted around my certainties. I spent the night computing them, point by point, my pencil scratching until the map was complete.

Figure 1: Gaussian Process (Noiseless)




Figure 1: **The Fireside Map — A Gaussian Process Prior.** Contour map of the imagined village. Solid circles mark the landmarks described by the villager, treated as noiseless. Solid lines are the GP posterior mean $\boldsymbol{\mu}_*$. Shaded regions show the variance $\sigma_*^2 = \text{diag}(\Sigma_*)$. Tight near observations, ballooning outward.

This is the essence of Gaussian Processes: we infer a function by combining prior beliefs (via the kernel) with observed data, assuming the observations are exact. I folded that map and slept, dreaming in Gaussians.

2. The Road Disagrees

I arrived at the valley on a Tuesday of thin, watery sunlight. The path, which my map told me would bend left beside a great oak, bent south instead. The oak was there,

gnarled and ancient, but it stood thirty paces east of where I had drawn it. I frowned, corrected my mental coordinates, and walked on.

The pond I sought—the kidney bean—was a meadow of yellow grass, dry and crackling underfoot. The schoolhouse, when I finally found it, crouched not on the hillcrest but slumped in a hollow, half hidden by brambles. Nothing was where it should have been. The old man’s words had been a rough sketch, smeared by the years and the ale and the frailties of human memory.

I sat on a stone and stared at my useless, beautiful map. My GP had been a fair-weather friend. It had believed, with innocent purity, that the villager’s every utterance was exact. It had no language for “*I’m not quite sure*” or “*the pond might be a bit farther than I recall.*” It could only treat every landmark as a nail hammered into the fabric of truth.

What I needed was a map that could accept the tremble in a storyteller’s hand.

And so I took out a second set of tools, a philosophy I had learned but often neglected. Its name was **Stochastic Kriging**, SK.

SK is GP’s wiser, more weathered sibling. It, too, believes in smoothness, but it adds a crucial caveat: every observation comes with noise. The villager’s description of the oak’s position was not a perfect coordinate; it was a distribution, a cloud of probability centered roughly around his pointed finger. Even my own measurements, taken with a compass and a knotted rope, carried their own shiver of error.

In SK, I gave each landmark a *variance*. The oak was at this spot, I recorded, but I wrote beside it: ± 15 paces, because the man had said “just where the path bends,” and paths can shift. The schoolhouse’s location I marked with even wider uncertainty, because I had found it in a hollow I didn’t expect. I was no longer pinning down truths; I was recording estimates, each with its own measure of trustworthiness.

The mathematics changes in a small but powerful way. The observation model becomes:

$$y_i = f(x_i) + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, \tau_i^2) \quad (7)$$

The noise variance τ_i^2 is *not* the same for all points; it can differ. This is crucial. For the oak, which I trusted more after finding it, I set $\tau^2 = 100$ (paces squared). For the pond, relocated only by a vague gesture, I set $\tau^2 = 900$.

The prior over f remains a GP with kernel $k(x, x')$. The joint distribution now includes the noise. The covariance of the observed outputs $\mathbf{y}$ becomes:

$$\text{Cov}[\mathbf{y}] = K(X, X) + \text{diag}(\tau_1^2, \tau_2^2, \dots, \tau_n^2) \quad (8)$$

Let $\Sigma_n = \text{diag}(\tau_1^2, \dots, \tau_n^2)$. Then the posterior predictions for new points are:

$$\boldsymbol{\mu}_* = K(X_*, X) (K(X, X) + \Sigma_n)^{-1} \mathbf{y} \quad (9)$$

$$\Sigma_* = K(X_*, X_*) - K(X_*, X) (K(X, X) + \Sigma_n)^{-1} K(X, X_*) \quad (10)$$

This small change—adding the noise matrix—completely altered the map. The mean no longer had to pass exactly through the observations; it could smooth over them if the noise was high. The variances Σ_* now included the effect of that noise: even at an observed point, the predicted variance was non-zero, because the true $f(x_i)$ is uncertain due to the noisy measurement.

I re-labelled every landmark with its estimated noise variance, recomputed the posterior, and redrew the map. The contours shifted. The pale halos of uncertainty no longer shrank to zero at the points; they stopped at a floor equal to the noise. The map had learned humility.

Figure 2: Stochastic Kriging (Noisy Observations)




Figure 2: **The Corrected Map — Stochastic Kriging.** The same village, but each observation now carries an error bar (the size of τ_i). The posterior mean no longer interpolates exactly; it smoothes. The variance field shows non-zero uncertainty at the observation locations, especially where τ_i is large.

Stochastic Kriging extends GP by acknowledging that observations are noisy, replacing exact interpolation with probabilistic smoothing via a noise covariance term. I walked on, and for the first time the map and the land began a dialogue.

3. The Whispering Slopes

On the third day, I decided to climb the eastern ridge.

The map I had—now a sensible SK document—showed a gradual, unremarkable rise. But as I ascended, the ground beneath my boots told a different story. The slope was steep, far steeper than my contour lines implied. Midway up, I stopped, breathing hard, and placed my hand on the hillside. I could feel the angle of it, the insistent downward tug, the way the earth fell away into a hidden valley I hadn’t even suspected.

I had more than a point measurement now. I had a *gradient*.

In cartography, we often ignore the slope. We mark the height and move on. But the slope is rich with information. Knowing that at this spot the land tilts sharply eastward tells me that the hill crest must lie somewhere west, and the valley floor somewhere east. A single point with a gradient can constrain a whole region.

This is the idea behind **Gradient-Enhanced Gaussian Process**, GEGP. It is an extension of GP that doesn’t just record the height—it records the derivative, the tilt. Mathematically, if a standard GP says “nearby points are similar,” GEGP says “nearby points are similar, and the rate of change is also similar.” It builds a map that honors not only the spot heights but also the steepness of the trail.

I took out my inclinometer, a crude thing of wood and brass, and I measured the slope where I stood: 18° , pitching east-northeast. I marked the spot on my map—a

single new observation—but I wrote not just the elevation, 247 metres, but also the partial derivatives: $\partial f/\partial x$ and $\partial f/\partial y$. These were the whispers of the hill.

For GEGP, the observation vector now includes function values and partial derivatives. Let

$$\nabla f = \left[\frac{\partial f}{\partial x^{(1)}}, \dots, \frac{\partial f}{\partial x^{(d)}} \right]^\top. \quad (11)$$

The GP is still over the latent function $f(x)$, but we observe a mix. The covariance between any two derivative observations can be obtained by differentiating the original kernel $k(x, x')$:

$$\text{Cov}\left(\frac{\partial f(x)}{\partial x_i}, \frac{\partial f(x')}{\partial x'_j}\right) = \frac{\partial^2}{\partial x_i \partial x'_j} k(x, x') \quad (12)$$

The covariance between a function value and a derivative is:

$$\text{Cov}\left(f(x), \frac{\partial f(x')}{\partial x'_j}\right) = \frac{\partial}{\partial x'_j} k(x, x') \quad (13)$$

For my squared exponential kernel, these derivatives are easy to compute. Let $r = \|x - x'\|$. Then:

$$\frac{\partial}{\partial x'_j} k(x, x') = \frac{(x_j - x'_j)}{\ell^2} k(x, x') \quad (14)$$

$$\frac{\partial^2}{\partial x_i \partial x'_j} k(x, x') = \left(\frac{\delta_{ij}}{\ell^2} - \frac{(x_i - x'_i)(x_j - x'_j)}{\ell^4} \right) k(x, x') \quad (15)$$

I constructed a large joint covariance matrix $\mathbf{K}$ over all observations (elevations and gradients), blocked appropriately. The training output vector $\mathbf{y}$ contained f -values and gradients. Assuming no noise for now (pure GEGP), the posterior formulas remain structurally the same:

$$\boldsymbol{\mu}_* = \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{y} \quad (16)$$

$$\Sigma_* = k(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}^{-1} \mathbf{k}_* \quad (17)$$

where $\mathbf{k}_*$ is the covariance between the new point x_* and all observations, including cross-covariances with gradients.

The inclusion of that single gradient observation collapsed my uncertainty along the direction of the slope. It was as if the hill itself had lent me its voice.

Gradient-enhanced GP incorporates derivative observations, which provide directional constraints that dramatically reduce uncertainty. On the ridgeline, I felt a cartographer’s deep, quiet thrill. The land was speaking to me in two languages now—height and slope—and I was learning to listen.

4. The Shaky Hands

But the world is not content to give us clean slopes.

Figure 3: Gradient-Enhanced GP




Figure 3: **The Power of a Gradient — GEGP vs. GP.** Left: Standard GP after one elevation measurement; symmetric uncertainty balloons. Right: GEGP after one elevation and one gradient; the variance is highly anisotropic, squashed along the slope direction, giving a much tighter map.

That afternoon, a rainstorm swept through the valley. The light turned grey and slippery. My hands were cold, and the inclinometer trembled. I took a gradient reading on the northern slope, but I knew it was uncertain. Had the angle been 22° , or 19° ? My numb fingers couldn't say.

I had entered a realm where even the gradients were noisy.

This is where pure GEGP falters. It assumes the gradient observations are perfect, just as naive GP assumes point observations are perfect. If you feed it a shaky gradient, it will over-trust that jitter, and the map may develop strange, unwarranted kinks.

The remedy is to marry the noise-handling of SK with the gradient-awareness of GEGP. The offspring of that union is **Gradient-Enhanced Stochastic Kriging**, GESK. It is the most complete of my map-making philosophies. GESK can accept a point measurement with uncertainty, and a gradient measurement with its own separate uncertainty. It weighs them all—height, noise in height, slope, noise in slope—in a single coherent model, producing a posterior map that accounts for all these trembling threads.

I recorded my rain-slick observations: elevation 312 m (± 4 m, due to the mist), slope 20° ($\pm 3^\circ$). In my notebook, I drew little error distributions next to each number. Then I invoked GESK.

Mathematically, I model noisy gradients as:

$$\widetilde{\frac{\partial f}{\partial x_i}} = \frac{\partial f(x)}{\partial x_i} + \epsilon_i^g, \quad \epsilon_i^g \sim \mathcal{N}(0, \sigma_{g,i}^2) \quad (18)$$

The full observation vector thus includes noisy heights and noisy gradients. The noise covariance matrix Σ_n is now diagonal, but with separate entries: for a function observation y_j , the noise variance is τ_j^2 ; for a gradient observation of the k -th partial derivative at some location, the noise variance is $\sigma_{g,jk}^2$. Then the joint covariance of all observations is:

$$\mathbf{K}_{\text{obs}} = \mathbf{K}_{\text{joint}} + \Sigma_n \quad (19)$$

where $\mathbf{K}_{\text{joint}}$ is the block matrix of covariances between all function and derivative evaluations (as in GEGP). The posterior formulas are identical in form:

$$\boldsymbol{\mu}_* = \mathbf{k}_*^\top \mathbf{K}_{\text{obs}}^{-1} \mathbf{y} \quad (20)$$

$$\Sigma_* = k(x_*, x_*) - \mathbf{k}_*^\top \mathbf{K}_{\text{obs}}^{-1} \mathbf{k}_* \quad (21)$$

Now $\mathbf{k}_*$ includes covariances with both function and gradient observations. The model naturally down-weights a noisy gradient observation relative to a clean one. The variance field becomes a rich tapestry: it is low near reliable points, stretched by clean gradients, but balloons where gradients are poorly known.

Figure 4: GESK (Noisy Gradients)




Figure 4: **The Full Toolkit — GESK.** A 2D map with circle markers (elevations, with error bars) and arrow markers (gradients, with fuzzy directional cones indicating gradient noise). The variance field shows suppression near clean gradients and elongation where gradient noise is high. The model fuses all information seamlessly.

GESK unifies noisy observations and noisy gradients into a single probabilistic framework, weighting each source of information by its reliability. I camped that night under a dripping canopy of leaves, and I understood that a map is never finished. It is a living document, a negotiation between what we have sensed and how much our senses can be trusted.

5. The Compass of Curiosity

With GESK in my hands, I had a powerful, noisy, slope-aware model of the valley. But I was not a passive collector. I was a traveler with a purpose and a finite number of boot-leather days. I could not visit every square foot of that village. I had to choose where to go next.

This is where **Active Learning** becomes the expedition leader. Active Learning is not a model; it is a strategy. It asks: *given my current map, which single new observation would most improve my knowledge?*

There are many possible answers, each a different path through the trees. The simplest criterion is to go where the map’s uncertainty is highest—the large, pale swaths of

variance. That is pure exploration, a walk toward the unknown. Another approach is to seek where the *gradient* is most uncertain, hoping that a well-placed slope measurement will lever open a whole region. Or one might balance exploration with the expected impact on the final map’s accuracy, computing an acquisition function that quantifies the value of each candidate point.

I sat with my GESK map and ran my finger over its pale patches. One such patch lay in the northwest, where the river’s course was a ghostly suggestion. Another lay on a steep south-facing slope where my sole gradient reading was rain-smeared and dubious. I decided to visit the south slope first: a single clear gradient there would anchor the entire hillside. The next day, I would trace the river.

Mathematically, let the current variance at a location x be $\sigma^2(x)$, taken from the diagonal of the posterior covariance Σ_* . A simple acquisition function is to pick the point of maximum variance:

$$x_{\text{next}} = \arg \max_{x \in \mathcal{X}_{\text{cand}}} \sigma^2(x) \quad (22)$$

This is *uncertainty sampling*. It sends me into the pale halos.

If I had a more global goal, such as reducing the total integrated prediction error over the entire domain, I could compute the expected reduction in variance at all points if I were to sample at a candidate x . The integrated mean squared error (IMSE) criterion sums the current variance over the region and subtracts the expected reduction. In practice, I used a grid of candidate locations and greedily picked the one that minimized the IMSE. For gradient observations, the acquisition function would favor locations where a gradient measurement would most slash the uncertainty. Crucially, these expected improvements can be computed analytically because the GP’s update depends only on the kernel and the new measurement location, not on the actual value.

Step by step, my compass of curiosity guided me. After each new measurement—each point recorded with its noise, each slope felt with its trembling—I updated the map. The pale halos shrank. The contours settled. The valley began to make sense, not as a perfect shape, but as a shape with known limits to my knowledge.



Figure 5: **The Active Learning Sequence.** Four panels. Top-left: initial GESK map with high uncertainty. A red ‘X’ marks the point of maximum variance (or highest IMSE reduction). Top-right: after adding an elevation there, uncertainty drops. Bottom-left: next iteration chooses a location where a gradient would help most. Bottom-right: final map with thin, even uncertainty. The sequence shows the map sharpening with each deliberate footstep.

Active Learning turns modeling into a decision process by selecting future samples to maximally reduce uncertainty or prediction error.

6. The Final Map, and the Lesson of the Land

On my last evening in the village, I sat on the porch of the actual schoolhouse—the one I’d found in the hollow—and I spread my final map on my knees. It was a thing of beauty and honesty. It had:

- The initial GP dream: smooth, trusting, ultimately false, with the elegant posterior $\mu_* = K_* K^{-1} \mathbf{y}$.
- The SK correction: noisy observations woven in with their proper weights, adding the noise matrix Σ_n so that the map became resilient to shaky stories.
- The GEGP revelation: slopes that spoke of unseen crests, using gradient covariances derived from $\partial k / \partial x$, dramatically reducing uncertainty with gradient whispers.
- The GESK synthesis: a full embrace of a world where both heights and slopes tremble, placing noise variances on derivatives and giving a map that respects trembling hands.
- The Active Learning journey: the deliberate, curious footsteps that, one by one, asked the land its secrets and turned pale ignorance into shaded confidence, guided by acquisition functions like $x_{\text{next}} = \arg \max \sigma^2(x)$.

I traced a finger along the final river line. It was not the ribbon the old man had described, but it was the true river, bounded by its uncertainty, honored in its reality.

Figure 6: Final Map




Figure 6: **The Final Map — A GESK Surface After Active Learning.** A rich contour map, dense with observations. Elevation markers show small error ellipses; gradient arrows cluster on slopes. The variance field is thin and uniform, a quiet hum of residual uncertainty. The legend notes the active learning path.

The old man had given me a story. GP had woven that story into a tapestry. SK had taught me that the threads were uneven. GEGP had let me feel the warp and weft of the ground itself. GESK had taught me to combine all of it, trembling hands and all. And Active Learning—Active Learning was simply the courage to ask the next question, to walk toward the pale spaces on the map, and to listen.

As I folded the map and prepared to leave the village, I realized that the journey had not been about a destination. It had been about learning how to see.